

RAMKISHORE NAMALA

Overland Park, KS | ramkishorenabila2002@gmail.com | +1 (720) 914-5016

[linkedin.com/in/ramkishore-namala](https://www.linkedin.com/in/ramkishore-namala) | github.com/ramkishore-namala

PROFESSIONAL SUMMARY

AI/ML Engineer with 4 years of experience designing and deploying production-grade machine learning systems, LLM-powered agents, and end-to-end RAG pipelines, backed by a Master's-level foundation in Computer Science. Proficient in Python, PyTorch, TensorFlow, LangChain, LangGraph, and LlamaIndex, with hands-on expertise in fine-tuning, NLP-to-SQL generation, semantic search, and MLOps tooling (Docker, Kubernetes, CI/CD) across AWS and Azure. Proven record of shipping scalable AI features that improve accuracy and enable direct, intelligent access to enterprise data, ready to contribute immediately to building customer-facing AI products.

TECHNICAL SKILLS

Programming Languages: Python, Java, C++, SQL

Machine Learning & Deep Learning: PyTorch, TensorFlow, Scikit-learn, Keras, Transformers, CNNs, Fine-Tuning

Generative AI & NLP: Large Language Models (LLMs), RAG, Fine-Tuning, Prompt Engineering, LangChain, LangGraph, LlamaIndex, Hugging Face, NLP-to-SQL, Agentic AI, Named Entity Recognition, Semantic Search, Text Classification

Vector Databases: Pinecone, FAISS, Weaviate, ChromaDB, Embeddings

MLOps & Deployment: Docker, Kubernetes (K8s), CI/CD, MLflow, REST APIs, FastAPI, Model Monitoring, Model Versioning

Cloud & AI Platforms: AWS (SageMaker, Lambda, S3), Microsoft Azure, Microsoft Fabric, Azure AI Foundry

Data & Databases: MySQL, MongoDB, Oracle, Pandas, NumPy, Spark

Core Concepts: Data Structures & Algorithms, Object-Oriented Programming (OOP), System Design, Distributed Systems

Methodologies: Agile, Scrum, Test-Driven Development, Cross-Functional Collaboration

PROFESSIONAL EXPERIENCE

Machine Learning Engineer – Applied AI

Jul 2022 – Jun 2024

OpenText | Hyderabad, India

- Enabled non-technical clients to query production data directly through natural language by integrating a Microsoft Fabric agent with Azure AI Foundry and deploying a semantic model that translates NLP prompts into SQL queries.
- Engineered an NLP-to-SQL prompt pipeline that converts user questions into precise, schema-aware SQL queries, allowing the system to retrieve exact data without manual query writing.
- Built an end-to-end RAG system powering an automation workflow agent, integrating LangChain, LangGraph, and LlamaIndex with vector databases to deliver context-grounded responses across enterprise document repositories.
- Improved document processing accuracy by 30% across enterprise content intelligence systems by building and optimizing ML pipelines using Python, Scikit-learn, PyTorch, and Transformer-based models for document classification and entity extraction.
- Developed LLM-powered applications using fine-tuning, embedding-based retrieval, and LLM APIs to strengthen intelligent document understanding and semantic search relevance.
- Exposed ML models through production REST APIs integrated into enterprise xECM platforms, enabling scalable, reliable backend inference for global clients.
- Implemented end-to-end MLOps workflows covering CI/CD automation, model monitoring, and version control, with containerized services deployed via Docker and Kubernetes (K8s) on AWS and Azure.

R&D; Engineer Intern

Jan 2022 – Jun 2022

OpenText | Hyderabad, India

- Prototyped and tested innovative AI/ML features in the enterprise content management domain, evaluating emerging technologies to inform future product roadmap decisions.
- Automated testing workflows for proof-of-concept modules, improving prototype reliability and shortening validation cycles.
- Collaborated with senior engineers to refine prototypes into scalable, maintainable designs ready for production integration.

Machine Learning Intern

May 2021 – Nov 2021

National Institute of Electronics & Information Technology (NIELIT) | India

- Built ML models for classification, clustering, and predictive analytics by training and evaluating algorithms in TensorFlow, PyTorch, and Scikit-learn on real-world datasets.
- Performed data preprocessing, feature engineering, and exploratory data analysis, and implemented deep learning architectures end-to-end to improve model performance.

PROJECTS

Stubble Management System | Java, Spring Boot, MySQL, React, REST APIs

- Won 1st place among 50+ teams at Hackathon 2021 (KL University) by designing and developing a full-stack platform connecting farmers and vendors for efficient stubble transactions.
- Enabled seamless management of user data and transactions by implementing secure REST APIs and MySQL database workflows behind a responsive React frontend.

Object Detection using CNN & Big Data | Python, TensorFlow, Keras, OpenCV, Hadoop

- Achieved 1st place at Project Expo 2021 by building a deep learning image recognition system capable of detecting face masks and objects using Convolutional Neural Networks (CNNs).
- Improved training efficiency on large-scale datasets by integrating Hadoop for distributed data processing.

EDUCATION

Master of Science in Computer Science

University of Central Missouri, USA

Bachelor of Technology in Computer Science

Koneru Lakshmaiah University, Hyderabad, India

CERTIFICATIONS

Cisco PCAP: Programming Essentials in Python | Aviatrix Certified Engineer – Multi-Cloud Networking Associate
Automation Anywhere – Advanced RPA Professional | Fundamentals of Data Analytics – NASSCOM